Understanding the Fundamentals of Vision Transformers (ViTs)

The ViT Architecture at a Glance

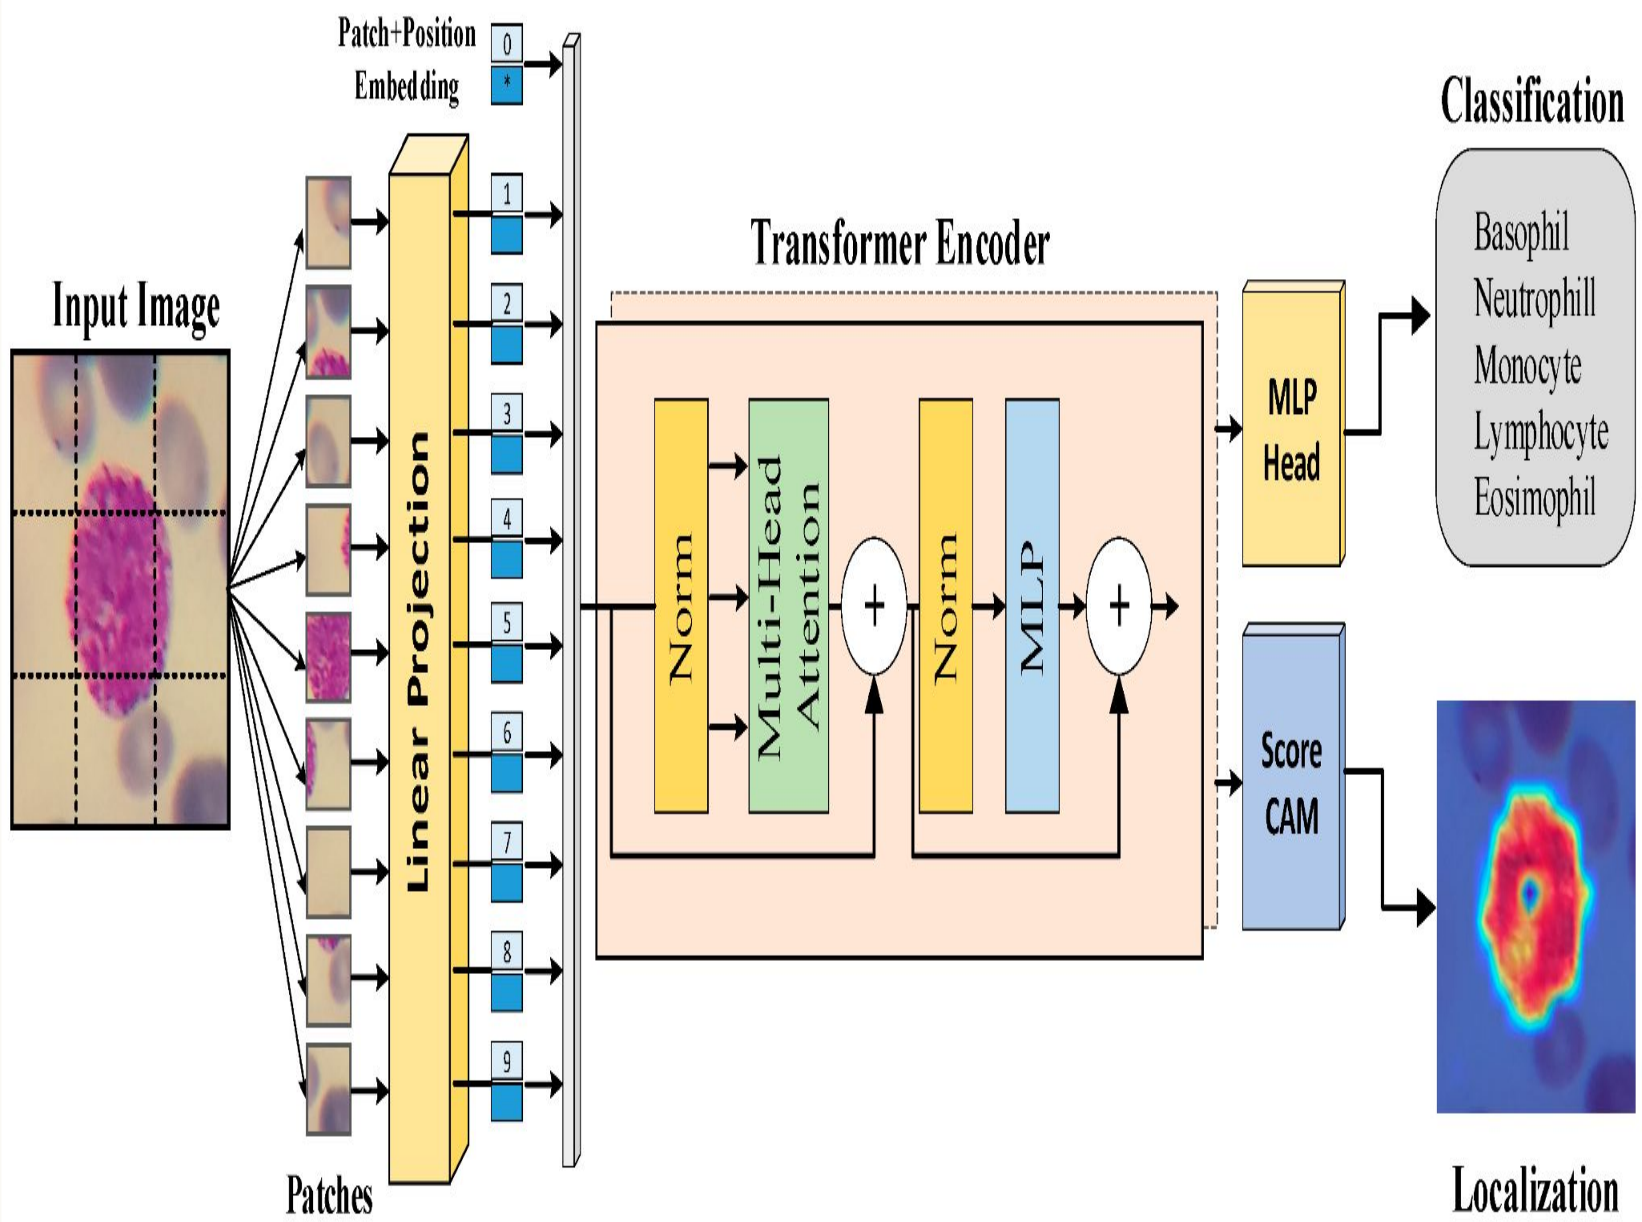

1 Patch Embedding

Breaking the image into fixed-size patches and converting them into linear embeddings.

2 Positional Encoding

Adding spatial information to the patch embeddings.

3 Transformer Encoder

The core attention mechanism learning relationships between patches.

4 Classification Head

A simple neural network for final prediction.

Step 1: Image to Patches

The first crucial step in a ViT is transforming a 2D image into a 1D sequence, similar to how text is handled in traditional Transformers. This is done by dividing the image into fixed-size, non-overlapping patches.

Image Partitioning: An image is split into a grid of smaller square regions (e.g., 16x16 pixels).

Linear Projection: Each 2D patch is flattened into a 1D vector and then linearly projected into a higher-dimensional embedding space.

Class Token: A special "class token" embedding is prepended to the sequence, similar to the CLS token in BERT. Its final state at the encoder output is used for classification.

Analogy: Think of a large painting cut into jigsaw puzzle pieces. Each piece is a patch.

Step 2: Positional Encoding

Unlike CNNs, which inherently understand spatial relationships due to their convolutional operations, Transformers are permutation-invariant. This means they don't naturally know the order or position of the patches.

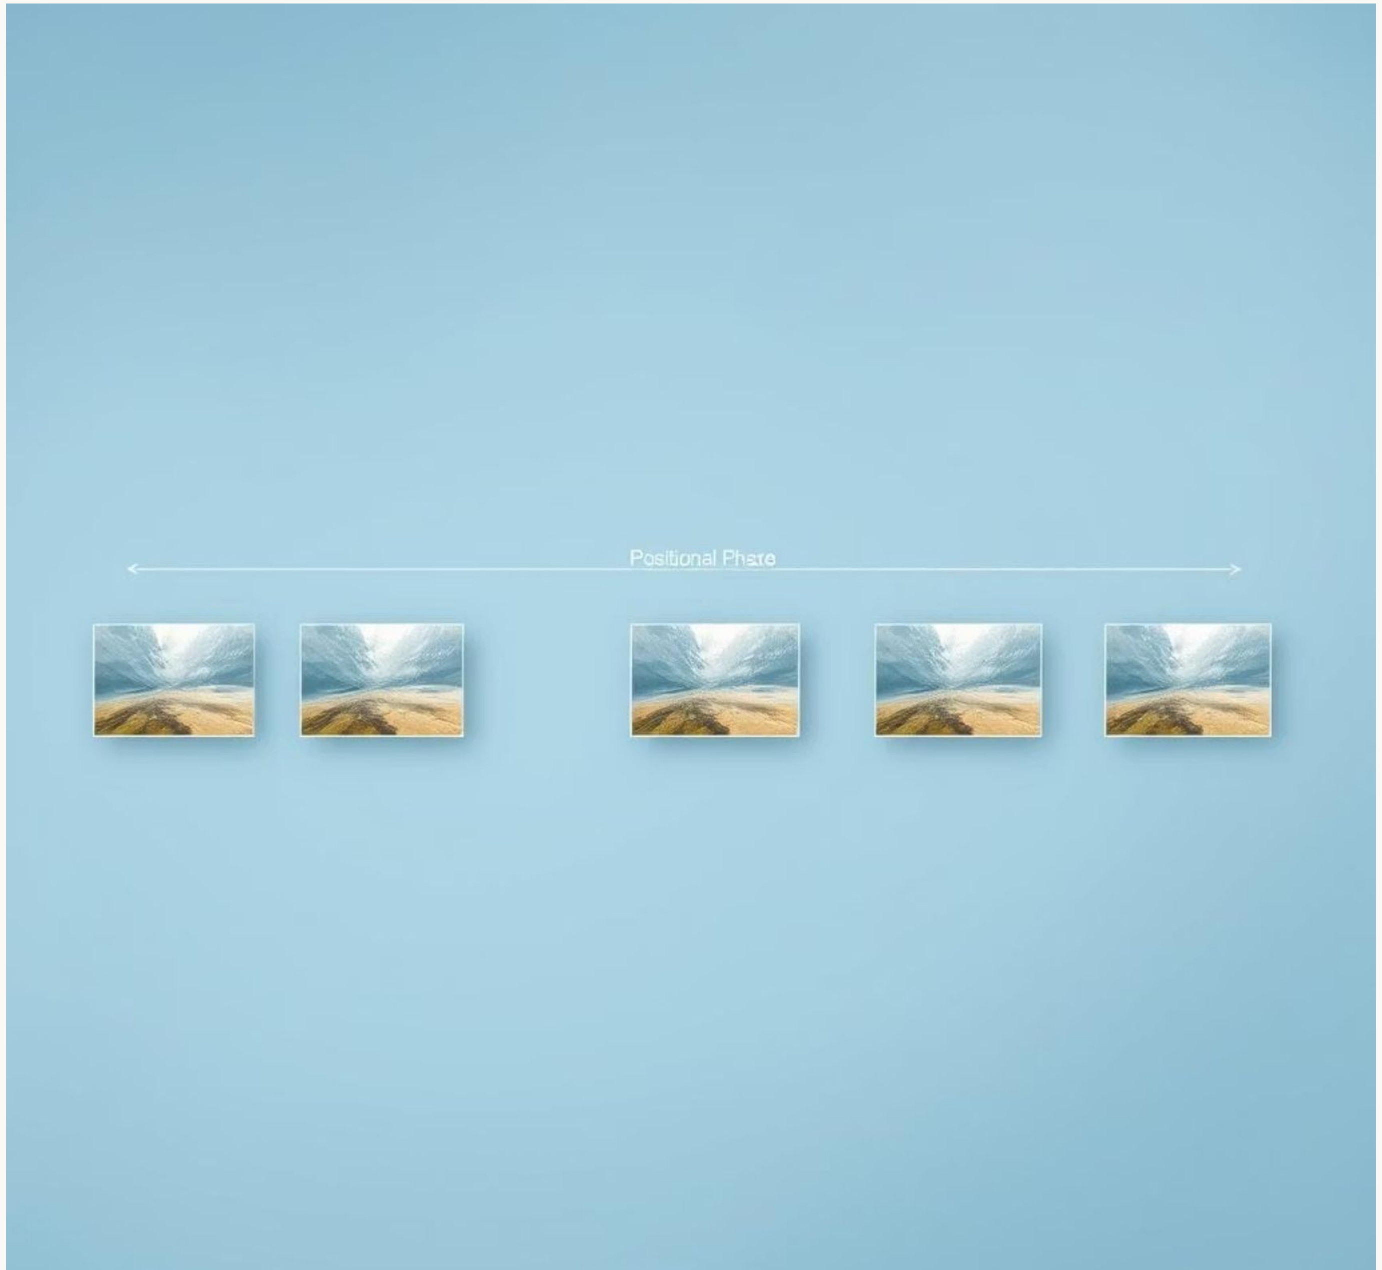

Adding Order: Positional encodings are added to the patch embeddings to inject spatial information.

Learnable or Fixed: These encodings can be fixed (e.g., sine and cosine functions) or learned during training.

Combined Input: The sum of the patch embedding and its corresponding positional encoding forms the final input sequence to the Transformer encoder.

Analogy: Giving each jigsaw puzzle piece a unique number so you know where it belongs in the original painting.

The Heart of ViT: Self-Attention

Self-attention allows each patch to "look" at all other patches in the image and weigh their importance when processing its own information. It's how the model captures global context.

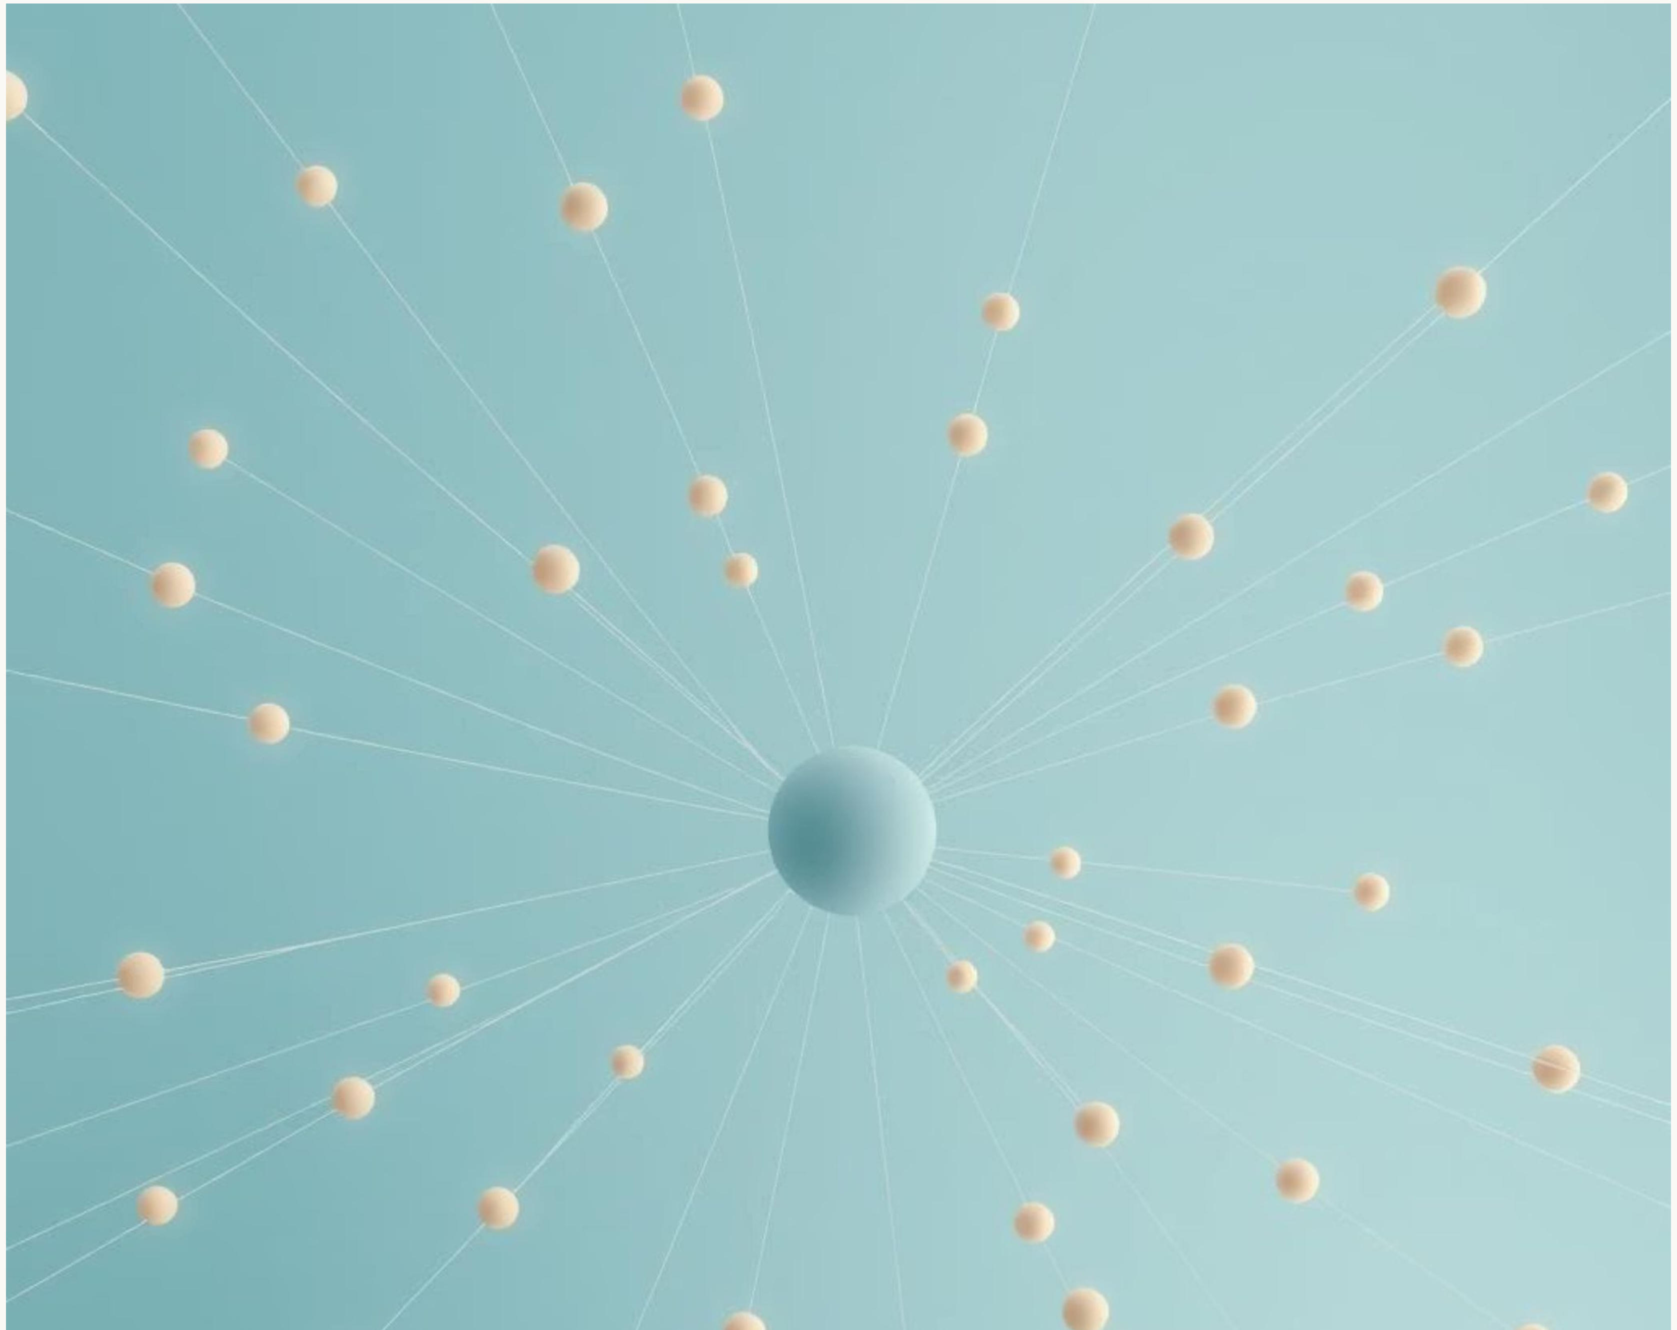

How it Works:

Query (Q): Represents the current patch.

Key (K): Represents all other patches.

Value (V): The actual content of other patches.

Attention Score: Calculated by the dot product of Query and Key, then scaled and passed through a softmax function to get weights.

Weighted Sum: These weights are then applied to the Value vectors to produce an output for each patch, enriched by context from all other patches.

ViT Pipeline: From Pixels to Prediction

Combining all components, the ViT processes an image through a sequence of transformations, ultimately producing a classification output.

Raw Image Input

The original image (e.g., 224x224 pixels).

Patching & Embedding

Image divided into fixed-size patches (e.g., 16x16), flattened, and linearly projected into embeddings. A class token is added.

Positional Encoding

Spatial information added to each patch embedding to retain positional context.

Transformer Encoder Blocks

Multiple layers of Multi-Head Self-Attention and MLP blocks process the sequence, learning global relationships.

Classification Head

The output of the class token from the final encoder block is fed into a Multi-Layer Perceptron (MLP) for final classification.

CNNs vs. ViTs: A Comparison

While both architectures excel in computer vision, their fundamental approaches lead to distinct advantages and disadvantages.

Feature	CNNs	ViTs
Locality vs. Global Context	Local receptive fields, hierarchy builds global view.	Global self-attention from Layer 1, direct long-range dependencies.
Parameter Count	Can be large, but often more efficient for smaller models.	Typically very large, especially for larger images/patches.
Data Efficiency	Inductive biases (locality, translation invariance) make them more data-efficient.	Requires massive datasets for pre-training to achieve comparable performance due to fewer inductive biases.
Computational Cost	Generally lower for inference on smaller inputs.	Quadratic complexity w.r.t. sequence length (number of patches).

Applications of Vision Transformers

ViTs have demonstrated remarkable performance across a wide range of computer vision tasks, often surpassing traditional CNNs, especially with sufficient data.

Image Classification

Identifying the main object or category within an image (e.g., animal, vehicle).

Object Detection

Locating and classifying multiple objects within an image with bounding boxes.

Semantic Segmentation

Pixel-level classification, assigning a class label to every pixel in an image.

Generative Models

Underpinning powerful image generation tools like DALL-E and Midjourney.